

Highlights

Domain-Validity-Gated Metamorphic Testing of Scientific ML Surrogates

Meng Li, Xiaohua Yang, Jie Liu, Shiyu Yan

- A domain-validity rubric screens candidate MRs for SciML surrogate testing.
- MR tolerances are gated by the measuring operator's floor ($O(h)$ for divergence).
- MR cards convert retained candidates into auditable executable test assets.
- A fail-closed claim ledger ties every reported number to a tracked artifact.
- K=6 rosters show where the admitted MR set bounds detector coverage.

Domain-Validity-Gated Metamorphic Testing of Scientific ML Surrogates

Meng Li^{a,b,c,1,*}, Xiaohua Yang^{a,b,c,2}, Jie Liu^{a,b,c,3}, Shiyu Yan^{a,b,c,4}

^a*School of Computing, University of South China, Hengyang, China*

^b*Hunan Engineering Research Center of Software Evaluation and Testing for Intellectual Equipment, Hengyang, China*

^c*CNNC Key Laboratory on High Trusted Computing, Hengyang, China*

Abstract

Context: Scientific machine-learning (SciML) surrogates approximate expensive simulations, but exact outputs for arbitrary inputs are unavailable (the oracle problem). Metamorphic testing addresses this through relations across executions, yet a candidate relation is not automatically valid: its preconditions, output mapping, and the scoring operator’s numerical floor determine whether a violation is meaningful.

Objective: We investigate how physics-derived metamorphic relations (MRs) can be transformed into auditable, executable, and interpretable oracle-free V&V assets for SciML surrogates, making the candidate-to-asset step explicit about the physical, numerical, and software conditions each relation requires.

Method: We propose (i) a domain-validity rubric admitting a candidate only when its tolerance dominates the operator’s numerical floor and its preconditions hold; (ii) an MR-card executable-asset format recording transformations, mappings, metrics, tolerances, exclusion rules, and relation-level verdicts; and (iii) a case-study protocol on MeshGraphNets cylinder-flow surrogates, with a claim ledger binding every result to an artifact.

*Corresponding author: Meng Li.

Email addresses: mlemon@usc.edu.cn (Meng Li), xiaohua1963@foxmail.com (Xiaohua Yang), jieliu@usc.edu.cn (Jie Liu), yanshiyu@usc.edu.cn (Shiyu Yan)

¹ORCID: Meng Li, 0000-0002-1074-1502.

²ORCID: Xiaohua Yang, 0000-0002-2977-1787.

³ORCID: Jie Liu, 0009-0008-1970-8347.

⁴ORCID: Shiyu Yan, 0000-0001-7626-5185.

Results: On a MeshGraphNets checkpoint, node permutation holds to machine precision, mirror-y is a bounded OOD-stress finding rather than an admitted symmetry, and absolute conservation stays deferred while a reference-relative guard passes; these readings recur unchanged across the same-family roster and a PhysicsNeMo replication. The same predicate applies unchanged to PINN and FNO surrogates, and on a second CFD task it rejects incompressible continuity on physical grounds rather than applying a fixed checklist. Seeded-fault and cross-SUT stress tests show how the admitted MR set bounds detector coverage in the studied settings, but these results are reported as an implication of the validity gate, not as a general fault-detection rate.

Conclusion: Validity-gated MR assets make oracle-free SciML V&V more auditable by recording why a relation is admitted, rejected, deferred, or interpreted as out-of-relation-domain. The study supports a reusable workflow for converting physics-derived relations into executable assets with typed verdicts. Broader generalization is future work.

Keywords: Metamorphic testing, metamorphic relation identification, oracle problem, scientific machine learning, neural surrogates, MeshGraphNets, domain-validity rubric

1. Introduction

Scientific machine-learning (SciML) surrogates are now used as software components for approximating expensive simulations. Mesh-based neural simulators such as MeshGraphNets predict physical fields on irregular meshes [1], but they also inherit a classical testing difficulty: for many candidate inputs there is no cheap exact output against which the program can be checked, the test oracle problem [2]. Held-out rollout error is therefore necessary, but it does not by itself tell a tester whether the software respects relations that should hold under physically meaningful transformations.

Metamorphic testing (MT) addresses the oracle problem by checking relations among multiple executions rather than a single expected output [3, 4]. Conservation laws, symmetry relations, nondimensional similarity, and continuity conditions can all suggest candidate metamorphic relations (MRs) for scientific software [5, 6, 7]. The testing problem considered in this paper begins after such candidates have been proposed. A physics-derived relation is not automatically a valid software test: mirror symmetry may depend on

geometry and boundary labels; a divergence relation may depend on a discrete operator, mesh weights, and boundary treatment; a transformed case may leave the relation’s domain before it reveals any SUT inconsistency.

For SciML surrogate software, this creates an interpretability problem for oracle-free V&V. When an MR check fails, the tester must ask whether the SUT violated a genuine relation, the relation was applied outside its validity conditions, the output mapping was incompatible with the software representation, or the scoring operator could not resolve the claimed effect. Treating all such outcomes as ordinary pass/fail MT verdicts makes the evidence hard to audit and easy to overclaim.

This paper therefore reframes SciML MR use as *validity-gated V&V asset construction*. It does not claim to invent MT, and it does not replace residual, uncertainty, accuracy, or trust-region diagnostics. Its contribution is to turn physics-derived candidate relations into auditable assets by combining measurement-floor admissibility, MR-card executable assets, and typed verdict ledgers. Physical knowledge, expert reasoning, LLM-assisted candidate lists, and NOETHER-style pattern organization [8] may suggest candidate relations, but the evidence supports bounded SciML V&V utility rather than general SciML reliability or real-world defect-detection rates.

Two principles organize the method. First, an MR is *admissible* only when its validity conditions hold for the SUT and the verdict is numerically decidable: the tolerance must dominate the intrinsic floor of the measuring operator, such as machine precision for an exact permutation relation, a mapping floor for an approximate geometric relation, or a discretization floor for a continuity relation. Second, relation-level verdicts are interpreted along two axes: relation-violation magnitude and domain-violation magnitude. This separates SUT inconsistency from out-of-relation-domain application, numerical-floor deferral, and inconclusive evidence. We refer to the resulting framework as domain-admissibility-gated, relation-indexed SciML V&V. The main empirical setting is MeshGraphNets-family cylinder flow, with the same admissibility predicate additionally exercised on a second CFD task and a second PDE family; these are bounded case studies, not cross-domain reliability estimates.

1.1. Research questions

The main research question is:

RQ0. How can physics-derived metamorphic relations for scientific-machine-learning surrogate software be transformed into auditable,

executable, and interpretable oracle-free V&V assets, while distinguishing genuine SUT inconsistency from relation invalidity and numerical measurement limits?

We decompose this into four questions:

- RQ1 – Admissibility.** What validity conditions determine whether a physics-derived candidate MR is admissible as an oracle-free test for a specific SciML surrogate SUT?
- RQ2 – Asset construction.** How can admissible MRs be represented as reusable MR cards and executable test assets with explicit transformations, mappings, metrics, tolerances, exclusions, and evidence records?
- RQ3 – Verdict interpretation.** How can relation-level verdicts separate SUT inconsistency from out-of-relation-domain application, numerical tolerance limits, and inconclusive evidence?
- RQ4 – Empirical utility.** In bounded SciML surrogate case studies, what evidence does the validity-gated workflow provide beyond roll-out accuracy, generic MR generation, and LLM-assisted candidate elicitation?

1.2. Contributions

This paper makes four scoped contributions. Table 1 records how each contribution is tied to a research question and to evidence reported later in the paper. The positive claim is methodological: a measurement-floor admissibility gate, a typed domain-inadmissibility verdict, and a seeded-fault diagnostic stress test whose MR-class-to-fault-class diagnostic is reported as stress-test evidence rather than a validated localization model.

Validity-gated V&V asset construction. The primary contribution is a workflow that transforms physics-derived candidate MRs into auditable executable testing assets. Each admitted MR is represented as an MR card and executable runner that records the transformation, mapping, metric, tolerance, exclusion rule, and ledger fields needed for independent inspection.

Measurement-floor admissibility. The method admits a relation only when its verdict tolerance dominates the intrinsic numerical floor of the measurement operator. For the deployed cylinder-flow mesh, the divergence floor is closed form, a leading-order predictor matches the measured floor

to within 0.5%, and a rigorous a-priori bound dominates it (Sec. 5.6); a general unstructured-mesh bound is left to future work.

Typed relation-level verdicts. The verdict scheme separates relation violation from domain violation, instantiating for physics-governed SciML the constraint-architecture pattern behind bug-versus-inapplicability reasoning [9, 10] and MetaTrimmer [11]. The added element is a typed classification of physical, geometric, boundary-condition, mapping, and operator-admissibility limits.

Bounded empirical utility. The case studies evaluate what this workflow adds beyond rollout accuracy, generic MR generation, and LLM-assisted MR candidate elicitation. Evidence is drawn from MeshGraphNets-family cylinder-flow runs, same-domain architectural variants, PhysicsNeMo replication, cross-family FNO/PINN checks, a second CFD task, and secondary baselines. The admissible MR set is also used to interpret fault-detection boundaries and blind spots, but this coverage reading is a discussion implication of the admitted relations rather than a separate formal research question.

2. Related Work

2.1. Mesh-based neural simulation

Mesh-based neural simulators learn dynamics on graph or mesh representations of physical systems. MeshGraphNets encodes mesh nodes and edges, propagates information by message passing, and predicts future physical fields through autoregressive rollout [1]. In this paper these models are treated as software systems under test (SUTs). The relevant testing question is therefore not only whether a rollout is accurate on held-out trajectories, but whether the SUT preserves declared relations under controlled transformations and explicit tolerances.

2.2. Metamorphic testing and the oracle problem

Metamorphic testing (MT) addresses programs lacking a per-input oracle [3, 4, 12]. Rather than checking each output against an expected value, MT checks a relation among source and follow-up executions. This framing has been applied to scientific software, PDE programs, simulations, and ML systems [13, 14, 5, 7], including recent applications to tabular ML such as credit scoring [15]. For SciML surrogates, however, an MR is not merely an input perturbation: its validity depends on governing equations, boundary

conditions, mesh representation, output mapping, and numerical tolerance. This makes MR admissibility part of the oracle construction problem itself.

2.3. MR identification for scientific and simulation software

Testing scientific software without a reliable oracle is a long-standing problem, and metamorphic testing is the most developed response in the surveyed literature [16]. MRs can be elicited from monotonicity, conservation, scaling, and domain-specific expectations [5, 6, 7, 17]. Predictive MR identification has been pursued with graph-kernel learning over scientific source code [18], by separating input and output patterns and by image-based output-pattern recognition for numerical solver codes [19, 20], and with large language models that discover metamorphic-relation variables for scientific software [21]. Symmetry-based MRs paired with statistical violation criteria have also been applied to deep trajectory predictors [22, 23]. Data-driven search can discover candidate transformations in complex physical software [24], and design assumptions can be encoded as MRs for CPS testing under control-plant invariants [25].

Separate work has made MR specification and applicability more explicit. General MT tooling has explored machine-readable MR representations, constraint definition, test-data-driven MR selection, and general MR specification languages [10, 26, 27]. These systems strengthen MR engineering in the large, but they do not by themselves decide whether a physics-derived relation is numerically measurable, in-domain for a transformed CFD case, or interpretable as a SciML V&V verdict.

These approaches are valuable candidate sources, but candidate identification is not the same as admissible SciML test construction. MR identification, including our own prior physics-derived MR construction for numerical solver and nuclear-power software [28, 19], is itself established work; this paper takes candidate MRs as input and contributes the admissibility gate rather than a new identification method. A failed SciML MR can mean that the SUT is inconsistent, that the relation was invalid under the boundary conditions, that the transformation left the intended physical regime, or that the tolerance was below the measuring operator’s floor. Fault-oriented MR work ranks relations by observed detection power [29, 30]; the present paper uses the admitted relation set later to interpret bounded coverage and blind spots, but treats that as a consequence of admissibility rather than as the primary novelty.

2.4. *Physics-based MT for learned scientific simulators*

Three contemporary strands are closest to this paper. Yang and Chui [31] and Reichert et al. [32] applied physics-derived MRs to hydrologic ML or conceptual models, showing that learned-model accuracy and relation consistency can diverge and that calibration or forcing regimes affect interpretation. Eniser et al. [33] introduced *relaxations*, numerical tolerances embedded in MR oracles, for stochastic RL policy testing. Duque-Torres et al. [9, 10, 26] and MetaTrimmer [11] address bug-versus-inapplicability reasoning with explicit MR constraints and automated constraint derivation.

The present work combines these ideas in a SciML-specific V&V workflow. From physics-based MR testing it takes relation-level checks; from relaxations it takes the need for calibrated tolerances; from violation-attribution work it takes the need to avoid treating inapplicability as a SUT failure. The added method requirement is that the admissibility decision, tolerance rationale, executable asset, and typed verdict are recorded together and tied to evidence artifacts.

2.5. *SciML V&V, residuals, UQ, and failure modes*

SciML reliability research has developed complementary tools: physics residuals, uncertainty quantification, conformal prediction, certified error bounds, and failure-mode analysis [34, 35, 36, 37, 38]. PINN training can exhibit gradient-flow pathologies that bias trained models away from the governing residual [39, 37], sharpening the need for checks that do not assume a converged surrogate. These diagnostics are not competitors to MT. A residual, equivariance error, or uncertainty score becomes an executable MR only when paired with a source case, follow-up transformation, tolerance rule, exclusion rule, and verdict interpretation. The proposed workflow uses such diagnostics as possible measurements inside a multi-execution oracle-free asset.

2.6. *Hybrid ML-solver trust regions*

Hybrid ML-solver frameworks use residual thresholds or error estimates to decide when a learned component should be trusted [40, 41]. Such systems operationalize runtime trust regions, whereas the present study acts offline: physically derived transformations are applied before deployment to estimate where relation violations occur.

The distinction from residual- and uncertainty-based trust estimation is deliberate. UQ, conformal prediction, residual-threshold trust regions, and

physical-consistency diagnostics that score a physically derived residual on observed outputs [42] locate unreliable behavior from observed inputs. The present method acts in relation space: it applies a controlled transformation and reports which necessary relation breaks under it, indexed by that transformation. The claim is complementarity, not superiority: relation-indexed evidence answers a different question from scalar accuracy, residual magnitude, or uncertainty.

2.7. What is new and what is not new

Metamorphic testing, MR identification, scientific-software MT, residual diagnostics, uncertainty quantification, LLM candidate generation, and NOETHER-style candidate organization are established or emerging sources of testing ideas. The contribution here is narrower: SciML MR use should be treated as a domain-validity and asset-construction problem. A candidate relation becomes a useful oracle-free V&V asset only after its physical basis, transformation preconditions, output mapping, tolerance, exclusion rule, executable artifact, and relation-level verdict are recorded. Recent work derives or refines domain-specific MRs for scientific programs directly from their numerical models [43, 44]; the wedge here is the learned surrogate evaluated out of distribution and the requirement that a relation’s tolerance dominate the measuring operator’s numerical floor, rather than testing a numerical solver against derived relations.

Table 2 summarizes the gap. Existing work supplies candidate sources, relaxed oracles, constraint-based inapplicability handling, and SciML diagnostics. The missing combination is a validity-gated, executable, and auditable workflow for physics-governed SciML surrogate software. The coverage discussion later in the paper is intentionally secondary: the admitted relation set helps explain what the test suite can and cannot detect in the studied settings, but the primary contribution is the construction and interpretation of the MR assets themselves.

3. Method

3.1. Overview

The method answers RQ1–RQ3 by turning a physics-derived candidate relation into an auditable V&V asset only after three checks have been made explicit: admissibility, asset construction, and verdict interpretation. Figure 1 summarizes the workflow. Candidate generation is deliberately separated

from validity judgment: expert reasoning, physical laws, NOETHER-style candidate organization, and LLM-assisted lists may propose relations, but none certifies that a relation is valid for a specific SUT.

Terminology and verdict types. A verdict is read on two axes (Sec. 3.5): a *relation-violation* magnitude, how far the SUT departs from the relation relative to its tolerance or numerical floor, and a *domain-violation* magnitude, whether the transformed case is in domain, near the boundary, or out of the relation’s domain. A candidate relation is accordingly *admitted* (a large in-domain violation then reads as SUT inconsistency), *rejected* on physical grounds, *deferred* when no tolerance dominates the operator floor, marked *out-of-relation-domain*, or left *inconclusive*. Four MR families recur in the case studies: node-permutation equivariance, mirror-y reflection symmetry, discrete-divergence continuity/conservation, and exact mirror-y on a constructed symmetric mesh.

3.2. Candidate relation sources

Candidate relations are read from the physical model, the computational model, and code-level representation contracts [45]. Physical-model candidates include symmetry, conservation, and nondimensional relations. Computational-model candidates include discrete operators and numerical constraints, such as a divergence measurement whose decidability depends on mesh weights and boundary treatment. Code-level candidates include representation contracts such as node permutation equivariance.

This classification is used only to locate the source of a candidate. The same relation may be valid at one level and invalid at another: mirror symmetry may be a physical relation for a symmetric domain but not for an asymmetric mesh with incompatible boundary labels; conservation may be a physical expectation but not an executable verdict when the discrete measuring operator has an unresolved floor. The admissibility gate records where the relation survives.

3.3. Admissibility gate

A candidate relation is an *admissible MR* only when four conditions hold together. First, it must have a physical or software basis. Second, its transformation preconditions must hold for the source and follow-up cases. Third, boundary conditions and output mappings must remain compatible after




Figure 1: Validity-gated V&V workflow. A candidate relation is admitted only after the gate records its basis, preconditions, mapping, and numerical decidability; admitted relations become MR cards, executable runners, and relation-level evidence ledgers.

transformation. Fourth, the verdict must be numerically decidable: the tolerance must dominate the intrinsic error floor of the measuring operator. None of these conditions is optional.

The gate yields four design-time decisions: admit as an executable MR; downgrade to OOD-stress when the relation is physically informative but outside its exact domain; reject when basis, preconditions, or mapping fail; or defer when missing numerical evidence prevents a meaningful verdict. For example, node permutation can be admitted as a representation MR when the adapter and inverse mapping are deterministic; mirror-y is conditional on compatible geometry, boundary labels, and vector-component mapping; absolute discrete divergence is deferred when the reference field and measuring operator floor cannot support an absolute verdict.

3.4. MR card and executable asset format

A retained MR is represented as an MR card. The card records the fields listed in Table 3: identifier, source category, physical or software basis, source-case schema, follow-up transformation, preconditions, boundary-condition compatibility, output mapping, metric, tolerance and provenance, expected verdict classes, exclusion rules, artifact schema, and interpretation limits. The card is the reusable design artifact; the executable runner is the operational artifact generated from it.

Each executable asset includes transformation code, runner configuration, metric computation, verdict logic, and artifact recording. Figure 2 shows the resulting data flow. The ledger is part of the method rather than only replication packaging: every relation-level claim must be traceable to the source run, follow-up run, mapping, metric output, floor evidence, exclusion decision, and verdict.

3.5. Relation-level verdicts

An MR execution can produce: **pass** (relation holds within tolerance); **fail** (violated within the valid relation domain); **skip** (required precondition missing); **out-of-relation-domain** (the transformed case is outside the relation domain); **numerical-tolerance issue** (the measured effect is dominated by threshold or resolution uncertainty); or **inconclusive** (insufficient artifact).

These verdicts form the two-dimensional space in Figure 3. One axis is relation-violation magnitude, typically expressed as a violation-to-floor or




Figure 2: MR-card-to-executable asset data flow. The MR card specifies the transformation, mapping, metric, tolerance, exclusion rule, and artifact schema; execution emits metric files, logs, and relation-level ledger entries.

violation-to-tolerance ratio. The other is domain-violation magnitude, meaning how far the transformed case lies outside the relation’s validity domain. Low domain violation with high relation violation is the region that may be read as SUT inconsistency; high domain violation is out-of-relation-domain; a violation inside the error floor is a numerical-tolerance issue. This decomposition is the method’s answer to RQ3.

In the present study, the domain-violation coordinate is operationalized per relation from committed measurements. For mirror-y, $D = m/(m + 1)$, where m is the worst reflected-node placement error in median-edge-length units. The synthetic symmetric mesh scores $D \approx 0$, whereas the real asymmetric evaluation mesh scores $D = 0.51$. Other MR classes use their own precondition measurements. These D values denote structural placement within each relation’s own validity domain: D is a per-relation normalized coordinate, not a cross-relation calibrated metric, and the values cannot be averaged or ranked across MR classes; they are read per relation, not as a calibrated boundary measurement. A cross-relation calibration would need to normalize each class by its own admissibility boundary, mapping every gate threshold to a common $D_{\text{cal}} = 1$; that calibration is future work.




Figure 3: Two-dimensional verdict space. The horizontal axis records relation violation relative to the relevant tolerance or floor; the vertical axis records whether the transformed case is in domain, near the boundary, or out of the relation domain. Only high relation violation inside the valid domain is read as SUT inconsistency.

Retained MRs are interpreted with the three-level scientific-computing MR classification of Yang et al. [45]: physical-model, computational-model, and code-model relations. This classification helps explain which layer a relation can observe. It is not a validated localization model; the seeded-

fault results later in the paper are a bounded stress test of this interpretation protocol.

4. Experimental Design

4.1. Evaluation logic

The experiments evaluate whether the method yields auditable, executable, and interpretable MR evidence under bounded conditions. They do not constitute a model benchmark and do not claim that the tested surrogates are generally reliable. The design separates four evidence tiers: primary cylinder-flow evidence, supporting cross-family checks, secondary candidate-generation and rollout comparators, and stress tests used to interpret coverage boundaries.

4.2. Subject systems and evidence tiers

Primary evidence comes from a MeshGraphNets-family cylinder-flow case study: a trained checkpoint, its K=6 checkpoint roster over three official held-out trajectories, same-domain S4/S5 wider/deeper variants, a PointMLP coordinate network, and the NVIDIA PhysicsNeMo **MeshGraphNet** production implementation evaluated on official DeepMind cylinder_flow TFRecords. The primary unit is a relation-level cell: MR card, SUT/checkpoint, trajectory or source case, transformed follow-up case, metric, and verdict.

Supporting evidence checks whether the same admissibility predicate can be applied outside the primary mesh-surrogate setting: a K=15 PINN roster and trained FNO-2D checkpoints over 2D Burgers and heat data. These runs support transfer of the gate and verdict logic, not cross-family pass-rate estimation.

Secondary evidence includes LLM-simulated expert MR design, generic MR-generation scope contrasts, LLM-assisted candidate generation, rollout-accuracy diagnostics, and external witness reruns from the read-only Minimum-MR-SubSet repository. These are contextual comparators and provenance checks, not defeated competitors.

Stress-test evidence uses seeded-fault and cross-program checks to interpret where the admitted MR set can and cannot observe faults. This tier supports the bounded coverage implication; it is not a real-world defect-detection-rate estimate.

4.3. *Baselines and comparators*

Rollout accuracy, residual/UQ diagnostics, expert-designed candidates, generic MR generators, and LLM-assisted candidates are used to contextualize RQ4. Rollout accuracy asks whether a prediction is close to observed held-out data; an MR asks whether a declared relation holds under a controlled transformation. LLM and generic MR outputs are treated as candidate sources; the admissibility gate, not the generator, decides whether a candidate becomes an executable MR.

4.4. *RQ-to-evidence map*

Table 4 maps the research questions to evaluation evidence, artifacts, units of analysis, and claim boundaries. It is a design table rather than a result table.

4.5. *Efficacy parameters and statistical plan*

Primary parameters are admissibility decision, executable MR rate, verdict distribution, violation magnitude relative to tolerance or floor, exclusion rate, and relation-level interpretation yield. Secondary parameters include inter-rater agreement for candidate screening, flakiness rate, complementarity with rollout accuracy, and stress-test detection by seeded-fault class.

The study reports inferential statistics only where the sampling structure supports them and labels descriptive counts otherwise. Binary detector outcomes over the unified fault catalogue carry Wilson 95% intervals per detector; continuous violation magnitudes use exact Wilcoxon signed-rank tests with a rank-based effect size (Cliff’s δ), where the design supports paired comparison; roster medians carry paired bootstrap intervals; and the operator-floor sweep reports a log-log slope with a 95% confidence interval. Repeated cells from one trajectory or one checkpoint family are reported as descriptive cell summaries, not independent population samples.

4.6. *Ethics, integrity, and reproducibility*

The study does not involve human subjects or personal data. All SUT versions, datasets, checkpoints, scripts, and run logs are recorded when licensing permits. Failed, skipped, out-of-relation-domain, numerical-tolerance, and inconclusive outcomes remain in the ledger. LLM use is restricted to candidate generation and evidence organization, not validity judgment. No candidate MR is described as valid without a passing gate decision and MR card; no OOD violation is described as a program fault without supporting verdict evidence.

5. Results Analysis and Discussion

This section reports the evidence by research question rather than by execution chronology. RQ1 is evaluated through gate decisions, task-specific admissibility, and the operator-floor calibration. RQ2 is evaluated through MR cards, executable runners, and the claim ledger. RQ3 is evaluated through typed verdicts that distinguish SUT inconsistency, out-of-relation-domain application, numerical deferral, and inconclusive evidence. RQ4 is evaluated through bounded SciML executions and secondary comparators. Fault-detection and cross-SUT coverage results are discussed only as a bounded implication of the admitted MR set.

5.1. Claim-to-evidence map

The compact claim-to-evidence map binds every claim to a tracked artifact using reviewer-facing claim names; the full table is given in Appendix A (Table A1), and the authoritative runtime mapping remains in `claim-ledger.yml`.

5.2. MR-card-to-verdict map

Table 5: MR-card-to-verdict map for RQ1–RQ3.

MR card	Rubric decision	Runtime verdict	Interpretation boundary
Node permutation equivariance	Retained as representation MR	pass sanity; relative L2 = 0.0	Supports one executable representation path; does not establish model reliability.
Mirror-y equivariance (asymmetric eval mesh)	Exact relation out-of-relation-domain; downgraded to approximate OOD-stress	S0 pilot: 10/10 fail; primary upgrade: 180/180 fail across K=6 x 3 trajectories x 10	Bounded within-family OOD-stress violation across held-out trajectories; not by itself exact symmetry, cross-SUT, or geometry-independent evidence.

MR card	Rubric decision	Runtime verdict	Interpretation boundary
Mirror-y equivariance (synthetic symmetric mesh)	Exact relation admissible (bijection verified, offset $< 10^{-12}$, type-match 1.0)	S0 pilot: fail, rel L2 1.10; primary upgrade: 18/18 fail across K=6 x 3 input seeds	Exact-symmetry failure where the relation is admissible; synthetic no-obstacle OOD meshes, not accuracy or cross-SUT evidence.
Discrete divergence / conservation	Absolute mass-conservation MR deferred; reference-relative diagnostic retained	S0 pilot inconclusive: reference-relative non-regression guard; primary upgrade: 162/162 pass across K=6 x 3 trajectories x 9	Reference-relative diagnostic only; the absolute conservation relation remains deferred.
MGN S4/S5 variant workflow	Node permutation admitted; mirror OOD/conservation diagnostic/exact-symmetry decisions recorded	2/2 node-permutation passes; 60/60 mirror OOD failures; 54/54 conservation-diagnostic passes; 6/6 exact-symmetry failures	Same-domain MGN variant evidence, not an external SUT family.
PointMLP cylinder workflow	Node permutation admitted; mirror OOD/conservation diagnostic/exact-symmetry decisions recorded	9/9 node-permutation passes; 10/10 mirror OOD failures; 9/9 conservation-diagnostic passes; 3/3 exact-symmetry failures	Different non-MGN cylinder SUT; not PhysicsNeMo/EchoWave or production CFD evidence.

MR card	Rubric decision	Runtime verdict	Interpretation boundary
PhysicsNeMo	Node	Node	Production-framework
MGN Object-A	permutation	permutation	artifact-chain smoke
smoke workflow	admitted; mirror	passes on the	evidence only; no
	OOD/	smoke subset	full-scale production
	conservation	(relative L2 0.0);	pass/fail rate and not
	diagnostic	rollout, mirror	external-aerodynamics
	decisions recorded	OOD, and	evidence.
		conservation	
		diagnostic ledgers	
		recorded	
FNO periodic	Translation	FNO primary	Full rubric-to-verdict
translation and	admitted;	workflow upgrade:	FNO evidence with raw
conservation	periodic discrete-	24/24 translation	source/follow-up outputs
	conservation MR	passes, 24/24	and per-case ledgers;
	admitted-with-	conservation	outside cylinder-flow and
	reference-floor;	failures, and 6/6	broad neural-operator
	Dirichlet	rejected Dirichlet	claims.
	translation	exact-MR	
	rejected	executions	

5.3. RQ1–RQ4 primary cylinder-flow evidence

All pilots run on one real trained MeshGraphNets cylinder-flow surrogate, using the read-only Minimum-MR-SubSet checkpoint recorded in the experiment ledger. They exercise three different rubric outcomes; raw outputs, manifests, and metric ledgers are committed (Table A1).

Representation MR. Node-permutation equivariance holds to machine precision (relative L2 = 0.0 at tolerance 1e-6). This is a structural property of message-passing and is reported as a pipeline-implementation sanity check rather than model capability or accuracy evidence.

Geometric MR. The rubric classifies exact mirror-y equivariance as out-of-relation-domain for this mesh (the reflection is non-bijective, the worst reflected-node mismatch is about one edge length, and the cylinder is off-centre by 7.2 mm) and downgrades it to an approximate nearest-neighbour OOD-stress probe scored against a same-space mapping-error floor. The probe failed on 10 of 10 recorded eval frames (median relative L2 0.737, median V/floor 3.96, floor range 3.02–5.55x), with no frame skipped or inconclu-

sive. This is a bounded within-SUT frame-level OOD-stress result: one SUT, one checkpoint, one MR, one eval trajectory. The frames are consecutive, not independent samples, so the exact binomial 95% interval for 10/10 is wide ($[0.69, 1.00]$). The mapping-error floor derives from the same geometric mismatch that triggered the downgrade, so V/floor alone cannot separate a model-level violation from an amplified geometric artifact on this mesh. The symmetric-mesh run below resolves that ambiguity.

Continuity MR. A P1 discrete-divergence operator yields a non-negligible divergence even for the ground-truth field on this coarse mesh, so an absolute divergence-free tolerance is not calibratable and the absolute mass-conservation relation stays deferred. As a reference-relative diagnostic, the surrogate’s predicted next-state divergence stays within about 0.4–0.8% of the reference on the recorded eval frames; the interior-only ratio confirms this is not only a boundary-imposition artifact. This asserts no absolute conservation. Two caveats bound the diagnostic. First, the reference divergence is not yet decomposed into discrete-operator error, solver projection artifact, or genuinely non-solenoidal training data, so if the reference field is itself materially non-solenoidal the reference-relative ratio is a non-regression guard rather than a conservation measurement. Second, the diagnostic uses a 50% regression threshold (ratio > 1.5) on two eval frames only, so “passes” means “does not regress conservation by more than 50% on those frames,” not “conserves mass.”

Rollout-accuracy diagnostic. On the same eval trajectory, the surrogate’s one-step next-state prediction error ($v_{\text{pred}} = v_t +$ denormalized predicted delta, the trainer’s own convention) has median relative L2 0.0216 (mean 0.044; min 0.0116, max 0.0788) over the nine recorded transitions. The per-transition error is bimodal, so we quote both statistics. The mirror-y OOD-stress violation (median 0.737) is about 34 times the median accuracy and 17x the mean. The quantities are dimensionless relative L2 values of different objects, so this is an order-of-magnitude gap rather than a precise factor; it is one-step, not a free-running rollout-stability result.

Exact mirror-y on a symmetric mesh. We built a provably symmetric synthetic structured channel mesh whose reflection is a verified bijection (node-type match 1.0, reflection offset $< 10^{-12}$, edge set invariant), so the predicate admits exact mirror-y. On one constructed input the surrogate fails this exact equivariance test (relative L2 1.10). The normalizer control changes the violation only from 1.1032 to 1.1014, so the violation is dominated by learned message-passing weights. The boundary is synthetic, no-

obstacle OOD geometry: the result confirms missing exact equivariance but is not a calibrated in-distribution magnitude.

Seeded-fault detection. We re-implemented, from the read-only Minimum-MR-SubSet witness taxonomy, an injected-fault catalogue of 10 pipeline faults across five fault classes (boundary-condition, mesh-adjacency, normalization-scale, temporal-rollout, physical-channel), and used the paper’s own MRs as detectors on the model’s predicted update. The continuity MR detected the two boundary-condition faults and the gross normalization fault (divergence ratio 3.8–10.6 vs the 1.5 threshold); the symmetry MR detected a physical-channel and a mesh-adjacency fault (violation rising 69–142% above baseline); node-permutation equivariance detected none, because these faults preserve node-relabeling invariance. 5 of 10 mutants were detected by at least one MR, and the detections separate by MR class (continuity to boundary/scale faults, symmetry to physical-channel/adjacency faults), a first bounded test of the interpretation protocol of Section 3.5, suggestive rather than a validation. Three boundaries: the detected faults are gross corruptions that any divergence- or symmetry-sensitive detector would catch; the edge-drop fault is a near-miss (32% mirror-y change vs the 50% threshold) and the boundary faults are invisible to mirror-y by construction; and the remaining undetected faults shift the absolute output without crossing the scored-quantity thresholds, delimiting where these MRs are structurally insensitive. Against a rollout-accuracy monitor on the same faults the two are complementary: the edge-permutation and channel-swap faults the symmetry MR catches leave the rollout error within $1.3\times$ of its in-distribution baseline (0.0216, under a $2\times$ regression threshold), so the MRs surface relation violations an accuracy monitor leaves within band, while the gross boundary/normalization faults trip both and no fault is caught by accuracy alone (claim C38). It is one SUT, one checkpoint, one injected-fault catalogue.

The remaining subsections carry the same evidence without adding another main-text table: the MGN roster tests denominator robustness, the operator-floor sweep calibrates numerical decidability, the fault catalogue bounds detector geometry, and the PINN/FNO executions test whether the gate transfers beyond the original MeshGraphNets workflow.

5.4. *RQ4 same-task replication across checkpoints, trajectories, and architectures*

The single-checkpoint evidence is replayed on a K=6 MeshGraphNets roster and three official DeepMind held-out test trajectories acquired through the Minimum-MR-SubSet loader. S0 reuses the pilot checkpoint, so this is a replication/variant check rather than six independent SUTs. Node permutation is exact on all checkpoints. The primary empirical scope upgrade gives a K=6 x 3 trajectories x 10 mirror-y OOD-stress grid with 180/180 failures, a K=6 x 3 trajectories x 9 conservation-transition grid with 162/162 reference-relative passes while absolute conservation remains deferred, and a K=6 x 3 exact-symmetric-mesh input grid with 18/18 failures. The upgrade removes the single-source-trajectory denominator while remaining within one architecture family and dataset, so per-cell counts are descriptive summaries, not independent-trial inference.

Same-task multi-architecture replication. Three further cylinder-flow executions test whether the verdict pattern is an artifact of one implementation. A same-domain S4/S5 wider/deeper MeshGraphNet variant package reproduces the pattern (2/2 node-permutation passes, 60/60 mirror OOD failures, 54/54 conservation-diagnostic passes, 6/6 exact-symmetry failures). A newly trained PointMLP coordinate network (a different architecture class with no message passing) reproduces it as well (9/9 node-permutation passes, 10/10 mirror OOD failures, 9/9 conservation-diagnostic passes, 3/3 exact-symmetry failures, median rollout relative L2 0.0298). Third, the NVIDIA PhysicsNeMo **MeshGraphNet** production implementation, trained and evaluated on official DeepMind cylinder_flow TFRecords, reproduces the node-permutation pass and the mirror-y OOD-stress failure within a production framework rather than a research codebase. The same admissibility predicate, applied unchanged, produced the same typed decisions on all three; this replicates the verdict pattern across architectures on one task and dataset, and asserts no cross-dataset rate.

5.5. *RQ1/RQ3 second-task gate discrimination: compressible airfoil flow*

The strongest test of a domain-validity gate is whether it produces the correct *different* verdict when the physics changes. We apply the identical four-condition predicate to a genuinely second CFD task on a second official dataset: the DeepMind **airfoil** benchmark (SU2-simulated compressible transonic flow, 5,233-node meshes with density, pressure, and velocity), distinct from incompressible cylinder flow in both physics and data. The same

official PhysicsNeMo MeshGraphNet (2.33M-parameter official configuration: hidden width 128, 15 processor steps) is trained for 15 epochs on 25 official airfoil train trajectories across a K=6 checkpoint roster and evaluated per-trajectory on 40 official test trajectories, yielding 240 (checkpoint, trajectory) pairs. The predicate yields a *different typed inadmissibility structure*, and the difference is the result: **node-permutation equivariance is admitted and exact** (240/240, relative L2 0.0, training-independent representation contract); **incompressible divergence-free continuity is rejected at the physical-basis gate** (condition i) because the density varies by a median factor of 2.13x across the field, so $\nabla \cdot u = 0$ is physically false here, the same relation that on cylinder flow passes conditions (i)–(iii) and is only *deferred* at condition (iv), so the gate assigns a categorically different verdict *type* for the physically correct reason on each task; **compressible unsteady mass conservation** $\partial_t \rho + \nabla \cdot (\rho u) = 0$ is the physically-correct replacement, its absolute discrete form deferred (uncalibratable operator floor) and a reference-relative diagnostic recorded (median residual ratio 1.00 across the K=6 roster); and **mirror-y chord symmetry is rejected at the boundary-precondition gate** (condition iii) because the SU2 trajectories carry a non-zero angle of attack. The one-step rollout error is high at this training budget (median relative L2 1.00, comparable to the signal magnitude) and is reported only as a training-state diagnostic; notably, none of the four typed verdicts depends on rollout quality: the admission rests on a representation contract and the rejections/deferral on physical and numerical reasoning, which is why the typed structure transfers and discriminates across tasks. This is a gate-discrimination result on an official-architecture but deliberately under-trained airfoil surrogate (240 checkpoint-by-trajectory cells), not an official-checkpoint, converged-model, or SOTA-accuracy airfoil result.

5.6. RQ1 operator-floor resolution sweep

The P1 discrete-divergence sweep on symmetric meshes, over nine resolutions, gives slope 0.984 with 95% CI [0.975, 0.992] and $R^2 = 0.9999$, matching $O(h)$. Beyond the slope, the *absolute* floor is closed-form for the concrete deployed-scale mesh: since the operator is exact on affine fields and the reference field is divergence-free, the measured floor equals the geometry-weighted second-order Taylor remainder exactly. A leading-order predictor from the analytic Hessian matches the measured floor to within 0.5% (1.337 vs 1.343), and a rigorous a-priori upper bound from the Hessian’s global spectral norm

dominates it at every resolution. The gate thus rests on a derived floor, not an empirical estimate, for this mesh (a general unstructured-mesh bound is future work). This is why absolute conservation remains deferred while the reference-relative non-regression guard is reported. A classical flux-form finite-volume operator instead admits a fully executable conservation verdict (claim C34, supporting artifact).




Figure 4: Operator-floor calibration for the P1 divergence measurement. The log-log sweep shows the $O(h)$ floor used to decide whether absolute conservation is numerically admissible in the cylinder-flow setting.

5.7. Coverage implication: fault-detection robustness across the multi-checkpoint roster

To check whether the seeded-fault detection result of Section 5.3 reproduces beyond a single checkpoint and to delimit where its detectors are structurally insensitive, the 10-mutant catalogue is replayed against the K=6 multi-checkpoint roster across five input-permutation seeds (30 trials per mutant per detector), and two predeclared severity dimensions are swept: NS_double_scale $s \in \{1.1, 1.25, 1.5, 2, 4\}$ and PC_zero_vy partial fraction $p \in \{0.25, 0.5, 0.75, 0.85, 0.9, 0.95, 0.99, 1.0\}$ (the canonical mutant is $p = 1.0$). Detection is decided by the same predeclared thresholds as Section 5.3 (node-permutation tolerance 10^{-5} , conservation ratio threshold 1.5,

mirror-y relative-change threshold 0.5). Wilson 95% CIs are computed across the (SUT, seed) cells; raw outputs are committed (Table A1).

Per-mutant cross-family detection (R1). Four mutants are detected on every cell (30/30, CI [0.89, 1.00]) and five on none (0/30, CI [0.00, 0.11]), reproducing the C10 by-class localization: continuity $\rightarrow$ boundary/scale faults, symmetry $\rightarrow$ physical-channel/adjacency faults, node-permutation $\rightarrow$ none (these preserve node-relabeling invariance). One mutant is configuration-sensitive: mirror-y detects MA_permute_edges on the seed-replica family (20/20 over S0–S3) but not on the wider/deeper S4/S5 variants (0/10), a within-family rate of 20/30 (CI [0.49, 0.81]) invisible to a single checkpoint.

NS_double_scale severity sweep (R2). At every multiplicative scale $s \in \{1.1, 1.25, 1.5, 2.0, 4.0\}$ the detection rate stays at 0/6 (Wilson CI [0.00, 0.39]). A constant output factor crosses none of the three thresholds: the reference-relative conservation ratio stays below 1.5, mirror-y moves under 50% (source and mirror scale together), and the linear operation preserves node-permutation equivariance. The catalogue at the current thresholds is therefore structurally insensitive to multiplicative output scaling across the swept grid, not only at the canonical $2\times$ fault.

PC_zero_vy partial-fraction sweep (R3), a knife-edge symmetry blind spot. The refined grid resolves *where* detection collapses. At every partial fraction up to $p = 0.99$ the detection rate is 6/6 across the roster (Wilson CI [0.61, 1.00]); it falls to 0/6 (CI [0.00, 0.39]) only at the exactly-uniform $p = 1.0$. The transition is a step rather than a gradual decline; even 1%-asymmetric zeroing is caught on every checkpoint. The responsible detector is *node-permutation*: partial zeroing masks v_y on a node-index-selected subset, which is not relabeling-invariant; the symmetry MR never fires in this sweep (0/6 at all fractions). At $p = 1.0$ the fault becomes simultaneously permutation-invariant and mirror-y-symmetric (uniform $v_y = 0 = -0$ under reflection) and escapes every geometric detector at once. The canonical fault therefore sits in a measure-zero blind region that any partial severity exposes, a within-family finding showing a fault can be invisible to a geometric MR suite precisely when it shares the suite’s invariances.

Aggregate reading and Phase-3 unified catalogue. The 5-of-10 union detection rate from C10 reproduces across S0–S3 (4/10 on S4/S5). By-class localization (continuity $\rightarrow$ boundary/scale; symmetry $\rightarrow$ physical-channel/adjacency) replicates across the cylinder checkpoint roster (its cross-SUT transfer is examined below), and the insensitivity regions (multiplicative output scaling, the PC_zero_vy / adversarial-mutant blind

subspace) are bounded with Wilson 95% CIs. This reflects the detectors’ coverage geometry: each MR scores a single invariant, so a fault is caught only when it perturbs a measured invariant and is structurally invisible when it preserves all of them; closing the blind region calls for additional MRs probing the uncovered directions, not more mutants of the same kind. Phase 3 additionally emits a 60-entry unified fault catalogue (Table A1): 10 executed canonical MGN mutants, 2 executed adversarial MGN mutants, 24 closed-form output-level PINN probes, and 24 closed-form output-level FNO probes. The precision/recall summary (descriptive catalogue rates over non-independent checkpoint-by-seed cells, not independent-item estimates) is node-permutation 1.00/1.00, conservation 1.00/0.81, and mirror-y 0.94/0.55; the same artifact reports Cliff’s δ and Wilcoxon signed-rank tests for the PINN MR-B ratios at $n = 15$ per PDE: Cliff’s $\delta = 0.73$ (large) and Wilcoxon $p = 0.0084$ for diffusion-vs-Burgers MR-B violation/floor ratios. This inter-PDE test is a physics-magnitude comparison (Dirichlet-zero Burgers BCs predict near-zero boundary residuals while Neumann-zero-flux diffusion BCs permit larger boundary residuals, so the magnitude difference is expected by construction) and is *not* a gate-reliability test. The valid reliability statistic is the per-PDE Wilson CI on the MR-B pass rate: Burgers 13/15 ([0.62, 0.96]) and diffusion 7/15 ([0.25, 0.70]). The PINN/FNO probes are not retrained mutant checkpoints.

Adversarial mutants (R4). Two adversarial mutants test whether the blind spot is a subspace, not a point. A1 is magnitude-driven, with node-permutation 0/6 and conservation 0/6; A2 escapes every detector. These results bound the detector set rather than establishing validated localization.

Bounded cross-SUT coverage implication. The seeded-fault stress tests suggest a bounded implication: coverage follows the admissible MR set in the studied settings. We test this on the primary-scale PhysicsNeMo airfoil (Section 5.5), with the same 10-mutant catalogue and predeclared thresholds, over five test trajectories (claim C36). The non-zero angle of attack makes mirror-y domain-inadmissible, so the gate excludes that relation and the mirror-only cylinder-flow fault coverage is absent on airfoil. The admissible relations keep their scoped readings: conservation still detects the gross normalization fault (NS_skip_denorm, all five trajectories) and node-permutation still detects none; the one localization shared across both SUTs is continuity $\rightarrow$ normalization-scale faults, whose MR is admissible on both. This supports a relation-level coverage interpretation, but it is not a real-world defect-rate estimate or a formal research question. We read it

as a proposed, falsifiable organizing principle, the validity-coverage duality: the same admissibility gate that decides whether an MR yields a meaningful verdict also fixes which faults that MR can detect; the cross-SUT removal of mirror-y above is one refutable prediction it makes, stated as a falsifiable hypothesis rather than a calibrated coverage law.

Falsifiable boundaries for the implication. The coverage implication would be weakened by coverage that survives exclusion of the only MR measuring the affected invariant, by a fault caught while perturbing no admitted measured invariant, or by a blind spot closed by adding same-kind mutants rather than new relation directions. In this evidence set, the cross-SUT air-foil run supports relation-specific exclusion, while the knife-edge sweep at $p = 1.0$ and the adversarial mutants delimit blind subspaces. The claim remains qualitative and fault-class level on two CFD SUTs, reported as an interpretive implication of RQ1–RQ4 rather than as a general theorem.

Cross-program breadth. The Minimum-MR-SubSet sibling is used only as an external breadth audit: seven program types are inspected read-only and five CPU-only SUTs are replayed through this paper’s gate and typed verdict (claims C39–C40). These checks test protocol portability, not per-program reliability, superiority, real-world defect rates, or new MR claims. Appendix Appendix B gives the full results.

5.8. RQ4 secondary baselines and external-scope audit

This subsection reports the Secondary baseline and external-scope audit: LLM baselines are secondary exploratory scope contrasts, not competitors. Expert-LLMs proposed 25 candidates (24 unique), but rubric voting retained 4 and found no novel retained MR family; generic templates admitted 3/13; a one-shot LLM baseline had low rater agreement (Fleiss kappa 0.077). These results support the need for an admissibility gate, not baseline superiority.

The Minimum-MR-SubSet audit/rerun is likewise an external witness: 70 real ABD instances, 20 true-fault-class rows, three SciML/PDE PASS witness rows, and one held-out cylinder-flow MGN PASS witness rerun. It adds artifact-backed scope context without adding a second primary architecture/dataset.

5.9. RQ4 Cross-family PINN extension

The cross-family executions test two claims the cylinder-flow study cannot test by itself: that the admissibility predicate is family-agnostic (the same four conditions, applied unchanged, produce typed decisions on pointwise

PINNs and spectral FNOs), and that the conservation MR class yields full executable verdicts wherever its floor is calibratable: the FNO periodic conservation MR fails 24/24 against a calibrated case-level floor and the PINN reference-relative conservation passes 30/30, while the MGN open-boundary absolute variant is the one member of the class the gate correctly refuses. The deferral on cylinder flow is the fail-closed gate discriminating rather than a missing verdict.

To check whether the workflow is MeshGraphNets-specific, we use a K=15 PINN roster: fifteen Burgers and fifteen heat seeds. **MR-A** remains *vacuous by construction* for a pointwise MLP. The two non-trivial MRs do the cross-family work: **MR-B** passes on 13/15 Burgers seeds (violation/floor ratio mean 0.712, bootstrap CI [0.583, 0.875]; Wilson pass-rate CI [0.62, 0.96]) and is mixed on heat with 7/15 passing (mean 1.495, CI [1.142, 1.897]; Wilson pass-rate CI [0.25, 0.70]), two seeds near the residual-floor boundary fail on Burgers, while the Neumann-zero-flux heat BCs predict larger boundary residuals and yield a lower pass rate by construction; **MR-C** passes on all 15/15 of both PDEs (Wilson CI [0.80, 1.00]), with Burgers mean ratio 1.007 (bootstrap CI [0.997, 1.017]) and heat mean ratio 1.006 (CI [0.988, 1.022]). This is a two-PDE PINN seed roster rather than a PINN-vs-MGN benchmark or a generality claim across PINN architectures, PDEs, or training regimes.

5.10. RQ4 FNO primary workflow upgrade

The **FNO primary workflow upgrade** converts the earlier FNO roster into a second trained primary execution. Six small FNO-2D checkpoints cover Burgers and heat. For each checkpoint and four held-out generated periodic finite-difference cases, the workflow records rubric decisions, source and follow-up tensors, mapped outputs, metric ledgers, and relation verdicts. Periodic integer translation is admitted and yields **24/24 translation passes** (maximum relative-L2 violation below 10^{-5}). The **periodic discrete-conservation MR** is admitted-with-reference-floor because the finite-difference target supplies a case-level channel-sum drift floor; the trained FNO outputs exceed that calibrated floor on **24/24 conservation failures**. The Dirichlet translation candidate is rejected for 6/6 SUTs because it changes the boundary-value problem and is not executed as an exact MR. This is **not only admissibility evidence**: it is a full rubric-to-verdict FNO execution with raw source/follow-up outputs and per-case ledgers, while remaining outside cylinder-flow evidence, performance benchmarking, reliability, and broad neural-operator generalization. Separately,

the FNO admissibility roster is extended to $K = 15$ seeds per PDE so the periodic-translation admission decision is adequately powered: it is admitted on 15/15 per PDE (Wilson 95% CI [0.80, 1.00]), with near-machine-precision violation magnitudes (Burgers mean 4.5×10^{-8} , heat 1.4×10^{-7}).

5.11. *Cross-RQ boundary of the evidence*

The canonical blocked list is narrowed but still active: the evidence does not support external-dataset or geometry-independent pass/fail rates, Aero-GraphNet/DoMINO results, comparative superiority over baselines, general or real-world fault-detection rates, validated localization, runtime, reliability, or model accuracy claims.

5.12. *Cross-RQ interpretation and discussion*

5.12.1. *Interpretation of the scoped evidence*

The value of the current study is not that every MR finds a new fault beyond rollout accuracy. Retained MRs provide relation-level evidence under explicit transformations; when a violation or deferral occurs, the MR card and verdict rule help distinguish model inconsistency, relation-domain boundary, numerical tolerance issue, and inconclusive evidence.

A natural objection is that the admissibility predicate merely re-describes the three initial outcomes. The added runs answer this: the symmetric-mesh run uses the predicate out-of-sample and fails an admitted relation; roll-out accuracy shows the mirror-y violation answers a different question from one-step error; and the primary empirical scope upgrade preserves these readings across a $K=6$ roster, three held-out trajectories, and larger denominators. The PINN roster, FNO primary workflow, and Minimum-MR-SubSet MGN/PINN reruns broaden applicability context while preserving the reliability and geometry-independent-rate boundary.

A second objection reads the deferred absolute-conservation relation as a gap in the end-to-end pipeline. The portfolio answers this directly: the conservation class produces full executable verdicts wherever its measurement floor is calibratable (FNO periodic conservation fails 24/24 against a calibrated floor; PINN reference-relative conservation passes 30/30), and the gate refuses execution precisely on the one variant where a verdict could not be separated from the operator’s own discretization error: MGN open-boundary absolute conservation.

The rubric averts two concrete misreadings: absolute discrete-divergence would have read as a conservation *pass* despite reference-field divergence, and

asymmetric-mesh mirror-y would have read as a model fault rather than a geometric artifact.

5.12.2. *Boundary of claims*

The paper avoids four overclaims: that MRs always beat rollout accuracy; that NOETHER proves the physical validity of cylinder-flow MRs; that LLMs can automatically identify valid MRs; and that one MeshGraphNets family generalizes to all SciML surrogates. The safer claim is that a domain-validity-aware workflow makes MR identification and execution more auditable for a concrete class of SciML SUTs.

5.12.3. *Implications for SciML testing*

The scoped evidence shows how SciML testing can move from implicit expert checks to explicit MR assets that complement residuals, UQ, and accuracy by making transformations and verdict rules inspectable. The workflow targets a relation-indexed applicability map in relation space rather than residual space. We do not claim a completed applicability map in this paper; the present evidence is one bounded within-SUT point, and calibrated maps across SUTs, trajectories, and domain-violation scores remain future work.

6. Threats to Validity

We classify threats following Verdecchia et al. [46] and report against the Empirical Standards for software engineering research [47].

Construct validity. The main construct is not “model correctness” but validity-gated MR asset construction. A wrong physical premise, boundary-condition reading, output mapping, discrete operator, or tolerance could therefore change an admit/reject/defer decision. We mitigate this threat by requiring an MR card, explicit preconditions, typed verdicts, and an operator-floor argument before a relation is treated as executable evidence. The mitigation is still incomplete: the P1 divergence floor is derived for the deployed symmetric-mesh setting, not for arbitrary unstructured meshes, and the reference-relative conservation diagnostic remains a non-regression guard rather than an absolute conservation measurement. The seeded-fault catalogue also measures detector sensitivity to author-implemented mutants, so its rates do not estimate real-world defect prevalence or effectiveness.

Internal validity. SUT setup, checkpoint selection, random seeds, mesh preprocessing, normalization, and runtime nondeterminism may affect verdicts. The study records manifests, ledgers, and raw outputs, and it repeats

the primary cylinder-flow checks across $K=6$ checkpoints, three held-out trajectories, same-task variants, PointMLP, and PhysicsNeMo. These repetitions reduce the risk that the pilot is a single-run artifact, but they do not make the cells independent population samples. For airfoil, the surrogate is deliberately under-trained at the local budget; its evidence is read for typed gate discrimination, not for model quality.

External validity. The evidence covers incompressible cylinder flow, compressible airfoil flow, bounded PINN/FNO Burgers and heat executions, and a read-only/executed cross-program breadth check from the Minimum-MR-SubSet sibling artifacts. This is enough to show that the gate is not a one-off cylinder-flow checklist, but it remains far from a general claim about SciML surrogates. Broader neural operators, mesh simulators, production CFD solvers, PINN architectures, datasets, geometries, and training regimes remain outside the evidence. The cross-program breadth check reuses the sibling’s domain assets and mutant sets, so it supports only a qualitative coverage-geometry reading, not per-program reliability, new MR discovery, superiority, or real-world fault-detection rates.

Baseline and comparator validity. Rollout accuracy, residual/UQ-style diagnostics, generic MR templates, and LLM-assisted candidate generation are used as scope contrasts rather than defeated competitors. The LLM baselines depend on prompt design, model versions, and simulated raters; they are useful for showing why an admissibility gate is needed, but they are not a human-expert study and do not establish superiority over alternative MR-identification methods.

Conclusion validity. The study combines descriptive counts, Wilson intervals, bootstrap intervals, and exact nonparametric tests where the sampling unit supports them. Many reported cells still share checkpoints, trajectories, or generated cases, so they are not independent population trials. The primary MGN scope upgrade improves the weakest denominators within the stated scope, but the correct conclusion is bounded empirical utility for RQ4, not a calibrated reliability or coverage rate.

Reproducibility. Some workflows depend on old runtimes, public benchmark loaders, or non-redistributable checkpoints. The replication package therefore emphasizes scripts, manifests, run logs, claim ledgers, and available outputs. Reproducibility is strongest for artifact inspection and rerunnable CPU workflows, and weaker for exact recreation of every trained checkpoint environment.

7. Future Work

Future work should generalize the numerical admissibility gate beyond the deployed P1 symmetric-mesh setting to unstructured meshes, boundary treatments, and flux-form operators.

Second, the domain-violation axis should be calibrated across MR classes. Here, domain distance is a per-relation interpretive coordinate, not a cross-MR score. A calibration protocol would normalize distances within relation families, define in-domain/near-boundary/out-of-relation-domain anchor cases, build relation-family calibration sets, and use independent domain experts or raters to validate categories, agreement, and residual comparability limits.

Third, the MR-card workflow should be evaluated on production-scale SUTs and datasets, including external aerodynamics, production graph simulators, longer rollouts, and multi-physics pipelines.

Fourth, the coverage implication should be tested with independent real-defect corpora and independently authored mutants, asking whether admitted MR sets predict externally collected defect classes.

Finally, MR-card authoring, precondition checking, and ledger generation need automation, while validity decisions remain tied to physical, numerical, and software evidence rather than model consensus.

8. Conclusion

This paper asked how physics-derived MRs for SciML surrogate software can be transformed into auditable, executable, and interpretable oracle-free V&V assets. The answer is a validity-gated workflow: a candidate relation is first checked for physical basis, domain preconditions, representation mapping, and numerical decidability; admitted relations are recorded as MR cards and executable assets; and outcomes are interpreted with typed verdicts rather than a single pass/fail label.

The empirical evidence is bounded but concrete. On cylinder-flow MeshGraphNets, the workflow separates a representation pass, an out-of-relation-domain mirror probe, a deferred absolute conservation relation, and a reference-relative diagnostic; the same readings are replayed across checkpoints, trajectories, and same-task architectures. On airfoil, PINN, and FNO subjects, the gate changes or preserves verdicts according to the relation’s physical and numerical conditions. Seeded-fault stress tests

further show that the admitted MR set helps explain detector blind spots in the studied settings, without implying general defect-detection rates.

The contribution is therefore methodological: physically meaningful SciML metamorphic tests require explicit validity conditions, executable assets, raw evidence records, and relation-level verdicts. This turns candidate MRs from informal expert intuitions into auditable V&V assets whose boundaries are visible to reviewers and practitioners.

Preprint

An earlier preprint version of this manuscript is available on arXiv at <https://arxiv.org/abs/2606.17529> with DOI <https://doi.org/10.48550/arXiv.2606.17529>.

CRedit authorship contribution statement

Meng Li: Conceptualization, Methodology, Software, Writing – original draft. **Xiaohua Yang:** Supervision, Formal analysis, Writing – review & editing. **Jie Liu:** Investigation, Validation. **Shiyu Yan:** Data curation, Visualization.

Declaration of competing interest

The authors declare no conflict of interest.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used generative AI assistants for code scaffolding, prose copy-editing, and reference cross-checking. After using these tools the authors reviewed and edited the content as needed and take full responsibility for the content of the publication. No generative AI was used to fabricate, alter, or invent experimental data, results, or citations; every numerical claim is backed by a tracked artifact under `research_assets/runs/` and a claim-ledger entry.

Acknowledgments

This work was supported by the National Natural Science Foundation of China (NSFC) General Program (grant no. 12575176), the Hunan Provincial Education Department Project, China (grant no. 202502000728), the Natural Science Foundation of Hunan Province, China (grant no. 2025JJ70193), and an industry-funded research project (grant no. 230KHX060001).

Data availability

The replication package contains the manuscript source, assets, ledger, and outputs, and is archived on Zenodo at <https://doi.org/10.5281/zenodo.20702952>. Minimum-MR-SubSet audit/rerun artifacts are under `research_assets/runs/` and cite commit `9ef862ec37335b4834d0a1fb38b4b613af702f34`. The MeshGraphNets checkpoints derive from the public DeepMind cylinder-flow benchmark; no proprietary data was used.

Table 1: Research questions, contributions, and evidence locations. Coverage is listed as a bounded implication, not as a formal research question.

Question	Contribution focus	Evidence used in the paper	Primary section
RQ0	Validity-gated V&V assets	End-to-end workflow, MR cards, runners, evidence ledger	Method; Results
RQ1	Admissibility	Domain-validity rubric, physical/software preconditions, operator-floor sweep	Method; Results
RQ2	Asset construction	MR-card schema, executable transformations, mappings, metrics, tolerances, exclusions	Method; Experimental Design
RQ3	Verdict interpretation	Relation-level verdict classes, two-axis verdict space, ledger records	Method; Results
RQ4	Empirical utility	Cylinder-flow runs, variants, PhysicsNeMo, FNO/PINN checks, baselines	Experimental Design; Results
Impl.	Coverage boundaries	Seeded-fault stress tests and cross-SUT coverage removal check	Results; Discussion

Table 2: Closest-prior capability matrix. The table positions the paper against capabilities needed by RQ0–RQ4; it does not assert that prior work should have solved the SciML-specific combination.

Approach	What it supplies	What remains missing for RQ0–RQ4	What this paper adds
Scientific-software MR identification	Candidate relations from domain rules, source-code learning, LLM variables, and data-driven search	Candidate validity, numerical decidability, executable asset format, and typed SciML verdicts	Treats these methods as candidate sources, not validity authorities
Reichert et al.	Physics-derived MRs and informal filtering in a learned hydrologic surrogate	Explicit admissibility predicate, operator-floor tolerance, reusable MR cards, and typed verdicts	Formalizes the candidate-to-verdict path for mesh-based SciML surrogates
Eniser et al.	Relaxed MR oracles with numerical thresholds	Measurement-operator floor and relation-domain interpretation for deterministic SciML surrogates	Grounds tolerance in the scoring operator’s own numerical floor
Duque-Torres / MetaTrimmer	Binary bug-versus-inapplicability separation with explicit MR constraints and automated constraint derivation	A typed inadmissibility verdict and operator-floor admissibility for physics-governed SciML	Replaces the binary skip/proceed pre-filter with a two-axis typed verdict tied to physical, geometric, boundary, and operator limits
Formal MR specification and constraint tooling	Machine-readable MR descriptions, constraint phases, and supporting systems	Physics-specific admissibility dimensions, measurement floors, and executable evidence ledgers for SciML	Uses specification ideas but binds each verdict to relation-domain and numerical-decidability evidence
This paper	Physics, expert, and LLM-assisted candidates plus artifact-backed execution	Bounded case-study evidence, not general SciML reliability or baseline superiority	Constructs MR cards, runners, ledgers, and typed verdicts under a four-condition gate

Table 3: Admissibility gate, MR-card fields, and verdict semantics. This is the method table for RQ1–RQ3.

Method element	Required record	Allowed interpretation	RQ
Basis	Equation, boundary condition, nondimensional law, representation property, numerical assumption, or implementation contract	Relation may be considered only for SUTs where this basis applies	RQ1
Preconditions	Source-case constraints, follow-up transformation, preserved quantities, and exclusion rule	Missing preconditions produce skip or out-of-relation-domain, not SUT failure	RQ1
Mapping compatibility	Boundary-label compatibility and output mapping from follow-up output back to source coordinates/fields	Incompatible mapping rejects or downgrades the candidate	RQ1
Numerical decidability	Metric, tolerance, tolerance provenance, and measurement-floor evidence	A verdict is meaningful only when tolerance dominates the operator floor	RQ1
MR card	Identifier, source category, transformation, mapping, metric, tolerance, exclusions, expected verdicts, and artifact schema	The admitted relation becomes a reusable executable V&V asset	RQ2
Evidence ledger	Source/follow-up run ids, metric outputs, floor evidence, exclusions, logs, and verdict	The reported claim is auditable from artifacts rather than prose alone	RQ2
Verdict semantics	Pass, fail, skip, out-of-relation-domain, numerical-tolerance issue, or inconclusive	Only fail inside the valid domain may be read as SUT inconsistency	RQ3

Table 4: Evaluation design by research question.

RQ	Evaluation question	Evidence tier and unit	Artifacts	Claim boundary
RQ0	Does the full workflow connect candidate, gate, executable asset, verdict, and evidence?	Primary workflow trace; MR-card cell	MR cards, runners, claim ledger	Asset construction only; no general reliability claim
RQ1	Which candidates are admissible, rejected, downgraded, or deferred?	Primary and supporting gates; candidate/MR	Rubric records, operator-floor sweep, exclusions	Physical validity is relation- and SUT-specific
RQ2	Are admitted relations executable and auditable?	Primary execution cells; source/follow-up pair	Transformations, runner configs, metric files, logs	Execution completeness, not complete automation of MR discovery
RQ3	Do verdicts separate SUT inconsistency from domain and numerical limits?	Verdict ledger; MR-by-case cell	Two-axis verdicts, metric/floor ratios, exclusion logs	Per-relation interpretation; no cross-relation calibrated domain score
RQ4	What evidence is added beyond rollout accuracy and candidate-generation comparators?	Primary, supporting, secondary, and stress-test tiers	Rollout diagnostics, LLM/generic baselines, witness reruns, seeded-fault ledgers	Complementarity and bounded utility; no baseline superiority
Impl.	What blind spots follow from the admitted MR set?	Stress-test tier; mutant/MR cell	Seeded-fault catalogue and cross-SUT removal check	Bounded coverage interpretation, not real-world detection rate

Appendix A. Claim-to-evidence map

Table A1: Compact claim-to-evidence map for the current manuscript.

Claim	Status	Evidence	Boundary
Domain-validity rubric	Method claim	Rubric asset and manuscript method section	Establishes an auditable screening rule; physical validity still depends on each relation's stated preconditions.
MR-card executable assets	Workflow claim	MR cards and validators	Some cards remain protocol assets pending matched SUT evidence.
Baseline admissibility contrast	Observed	Expert-MR, generic-MR, LLM-MR baselines and claim ledger	Quantifies an admissibility gap without ranking methods as competitors.
Minimum-MR-SubSet external-scope audit	Observed secondary audit	External-scope audit JSON	external witness evidence from 70 real rows and three SciML/PDE true-fault-class PASS_WITNESS rows; does not add new primary SUT executions to this paper.
Minimum-MR-SubSet primary rerun	Observed external runtime rerun	ABD witness rerun report	PASS_WITNESS, kstar = 6, four active true fault classes, max signature rank 2, collapse false; real held-out cylinder-flow MGN runtime, not a second architecture or dataset.
Minimum-MR-SubSet PINN primary reruns	Observed second-SUT/PDE reruns	ABD PINN witness rerun reports	trained Burgers2D and Diffusion2D PINN witnesses, PASS_WITNESS, kstar = 1 each, five active true fault classes each; one-seed witnesses, no cross-SUT rate.
Node-permutation sanity check	Observed pilot	Node-permutation metric ledger	One representation-contract path on one pilot case.
Conservation diagnostic	Diagnostic; absolute relation deferred	Conservation metric ledger and report	Reference-relative guard only; absolute conservation remains deferred.
Mirror-y OOD stress (PC6-mirror-y-ood-stress)	Observed bounded pilot	Mirror-y rate metric ledger and claim ledger	Failed on 10 of 10 recorded eval frames; not a reliability, accuracy, baseline, multi-SUT, exact-symmetry, or geometry-independent claim.
Primary empirical scope upgrade	Observed	Primary-scope report and per-trajectory/checkpoint manifests	K=6 x 3 trajectories x 10 mirror-y OOD-stress grid fails 180/180; K=6 x 3 trajectories x 9 conservation-transition grid passes 162/162; K=6 x 3 exact-symmetric-mesh input grid fails 18/18; not a single-source-trajectory estimate and not a cross-SUT rate.
LLM role boundary	Process boundary	Method and ethics sections	LLMs organize candidates; the rubric decides validity.
Rollout-accuracy diagnostic	Observed	Rollout-accuracy metric ledger	Same-SUT one-step median relative L2 0.0216; mirror-y is ~34x larger.
Exact mirror-y symmetric mesh	Observed	Symmetric-mesh metric ledger	Exact relation admissible and fails (rel L2 1.10) on a synthetic OOD mesh; binary equivariance evidence.

Claim	Status	Evidence	Boundary
PC10-seeded-fault-detection	Observed	Seeded-fault metric ledger	Detector stress test: MRs catch 5/10 injected mutants with class-specific response patterns.
Multicheckpoint replication	Observed	E1 aggregate and per-SUT manifests in multicheckpoint/	K=6 checkpoints of one MeshGraphNets family; within-family replication.
Operator-floor resolution	Observed	Operator-floor sweep report	Log-log slope 0.984, 95% CI [0.975, 0.992] (interior 0.989), $R^2 = 0.9999$ over nine resolutions; calibrates the admissibility predicate on one mesh family.
Fault-detection robustness	Observed	E3 robustness report; Phase-3 unified catalogue	30-trial Wilson CIs per MGN mutant plus a 60-entry unified fault catalogue (10 canonical MGN + 2 adversarial MGN + 24 closed-form PINN output-level probes + 24 closed-form FNO output-level probes), by-detector precision/recall with Wilson CIs, and Wilcoxon/Cliff effect-size tests; PINN/FNO probes are closed-form output-level probes.
S4/S5 variant primary workflow	Observed	Variant report and raw ledgers	Same-domain wider/deeper MGN variants: 2/2 node-permutation passes, 60/60 mirror OOD failures, 54/54 conservation-diagnostic passes, 6/6 exact-symmetric failures; not PhysicsNeMo/EchoWave or cross-dataset reliability evidence.
PointMLP cylinder primary workflow	Observed	PointMLP report, checkpoint, and raw ledgers	Newly trained non-MGN row-wise coordinate network: 9/9 node-permutation passes, 10/10 mirror OOD failures, 9/9 conservation-diagnostic passes, 3/3 exact-symmetric failures, median rollout rel L2 0.0298; not PhysicsNeMo/EchoWave or production CFD evidence.
PhysicsNeMo MGN Object-A smoke workflow	Observed smoke-subset production-framework execution	PhysicsNeMo smoke report, newly trained checkpoint, raw outputs, and metric ledgers	NVIDIA PhysicsNeMo MeshGraphNet executes on a first-record official DeepMind cylinder_flow smoke subset: node permutation passes (relative L2 0.0), mirror-y is OOD stress, and conservation is reference-relative diagnostic only; not full production-scale PhysicsNeMo, AeroGraphNet, or DoMINO evidence.
FNO primary workflow upgrade	Observed	FNO workflow report, per-SUT ledgers, and raw source/follow-up outputs	Six trained FNO-2D checkpoints over Burgers/heat: 24/24 translation passes, 24/24 conservation failures under a periodic discrete-conservation MR, and 6/6 Dirichlet-boundary translation rejections; not only admissibility evidence.

Appendix B. Cross-program breadth

Cross-program generalization. Reusing the committed Minimum-MR-SubSet kill matrices, the coverage-geometry reading reproduces across seven program types in three families: neural surrogates, classical numerical solvers (wave, point-kinetics, Burgers), and a production Monte-Carlo physics code (OpenMC). For each, fault coverage is the union of its admissible MR set’s per-class coverage; structural blind spots are present in six of the seven; and the class-to-class mapping is program-specific, extending the SUT-specificity finding from two CFD SUTs to seven program types. Detection rate tracks the admissible MR set rather than being uniform (minimal MR sets reach 10–20% with large structural blind spots, richer sets up to 100%), which is the coverage-generation structure the principle predicts. This is a read-only generalization check from committed evidence, not an end-to-end execution of this pipeline on each program (claim C39).

End-to-end cross-program execution. Beyond that read-only reuse, this paper’s pipeline is executed end-to-end on five CPU-only SUTs imported read-only from the same sibling: four classical solvers (heat parabolic, wave hyperbolic, point-kinetics stiff-ODE, Burgers conservation law) and the sibling’s OpenMC subject run with the real OpenMC Monte-Carlo k -eigenvalue solver (single-energy-group infinite medium, verified against the closed-form k_∞), applying the four-condition gate, the typed verdict, and the coverage map to each program’s own domain MRs and seeded mutants. The gate is no rubber stamp: across the five it admits 38 candidate relations and rejects 7 (two on wave, three on point-kinetics, two on OpenMC) at the output-mapping condition, where a one-sided monotonicity relation is not falsified by a scrambled-observable probe, and defers none. Each admitted relation holds on the correct program (the in-domain, consistent region of the typed verdict) while the mutant runs populate the SUT-inconsistency region; every program keeps structural blind spots once the sibling’s declared-equivalent mutants are set aside, and the operator-class-to-coverage mapping is program-specific. Detection tracks the admissible set rather than being uniform (full-set mutation score 0.75 on heat, 0.20 on the other three solvers, and 0.56 on OpenMC, for the sibling’s committed mutant sets). The MRs and mutants are the sibling’s; only the gate and typed verdict are newly executed here, so this asserts no per-program reliability, no superiority, no new MR contribution, and, for OpenMC, no continuous-energy or whole-core reactor-physics result (claim C40).

References

- [1] T. Pfaff, M. Fortunato, A. Sanchez-Gonzalez, P. W. Battaglia, Learning mesh-based simulation with graph networks, in: International Conference on Learning Representations, 2021. arXiv:2010.03409, doi:10.48550/arXiv.2010.03409.
- [2] E. T. Barr, M. Harman, P. McMinn, M. Shahbaz, S. Yoo, The oracle problem in software testing: A survey, *IEEE Transactions on Software Engineering* 41 (5) (2015) 507–525. doi:10.1109/TSE.2014.2372785.
- [3] T. Y. Chen, S. C. Cheung, S. M. Yiu, Metamorphic testing: A new approach for generating next test cases, Tech. Rep. HKUST-CS98-01, The Hong Kong University of Science and Technology (1998).
- [4] S. Segura, G. Fraser, A. B. Sanchez, A. Ruiz-Cortes, A survey on metamorphic testing, *IEEE Transactions on Software Engineering* 42 (9) (2016) 805–824. doi:10.1109/TSE.2016.2532875.
- [5] U. Kanewala, T. Y. Chen, Metamorphic testing: A simple yet effective approach for testing scientific software, *Computing in Science and Engineering* 21 (1) (2019) 66–72. doi:10.1109/MCSE.2018.2875368.
- [6] X. Lin, M. Simon, N. Niu, Exploratory metamorphic testing for scientific software, *Computing in Science and Engineering* 22 (2) (2020) 78–87. doi:10.1109/MCSE.2018.2880577.
- [7] M. Olsen, M. S. Raunak, Increasing validity of simulation models through metamorphic testing, *IEEE Transactions on Reliability* 68 (1) (2019) 91–108. doi:10.1109/TR.2018.2850315.
- [8] M. Li, X. Yang, J. Liu, S. Yan, Noether: A constructive framework for metamorphic pattern discovery from operator algebras (2026). arXiv:2605.17390, doi:10.48550/arXiv.2605.17390.
- [9] A. Duque-Torres, D. Pfahl, C. Klammer, S. Fischer, Bug or not bug? analysing the reasons behind metamorphic relation violations, in: Proceedings of the 2023 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER), 2023, pp. 905–912. doi:10.1109/SANER56733.2023.00109.

- [10] A. Duque-Torres, D. Pfahl, Towards a complete metamorphic testing pipeline, in: 2023 IEEE International Conference on Software Maintenance and Evolution (ICSME), 2023, pp. 606–610. arXiv:2310.00338, doi:10.1109/ICSME58846.2023.00081.
- [11] A. Duque-Torres, D. Pfahl, C. Klammer, S. Fischer, Exploring a test data-driven method for selecting and constraining metamorphic relations, in: 2023 49th Euromicro Conference on Software Engineering and Advanced Applications (SEAA), 2023, pp. 370–377. arXiv:2307.15522, doi:10.1109/SEAA60479.2023.00063.
- [12] T. Y. Chen, F.-C. Kuo, H. Liu, P.-L. Poon, D. Towey, T. H. Tse, Z. Q. Zhou, Metamorphic testing, *ACM Computing Surveys* 51 (1) (2018) 4:1–4:27. doi:10.1145/3143561.
- [13] T. Y. Chen, J. Feng, T. H. Tse, Metamorphic testing of programs on partial differential equations: A case study, in: Proceedings of the 26th Annual International Computer Software and Applications Conference, 2002, pp. 327–333. doi:10.1109/CMPSAC.2002.1045022.
- [14] X. Xie, J. W. K. Ho, C. Murphy, G. Kaiser, B. Xu, T. Y. Chen, Testing and validating machine learning classifiers by metamorphic testing, *Journal of Systems and Software* 84 (4) (2011) 544–558. doi:10.1016/j.jss.2010.11.920.
- [15] Z. Ying, A. G. Bellotti, J. L. Breeden, D. Towey, Metamorphic testing and exploration for machine learning credit score models, *Information and Software Technology* 188 (2025) 107903. doi:10.1016/j.infsof.2025.107903.
- [16] U. Kanewala, J. M. Bieman, Testing scientific software: A systematic literature review, *Information and Software Technology* 56 (10) (2014) 1219–1232. doi:10.1016/j.infsof.2014.05.006.
- [17] M. S. Raunak, M. M. Olsen, Metamorphic testing on the continuum of verification and validation of simulation models, in: 2021 IEEE/ACM 6th International Workshop on Metamorphic Testing, 2021, pp. 47–52. doi:10.1109/MET52542.2021.00015.
- [18] U. Kanewala, J. M. Bieman, A. Ben-Hur, Predicting metamorphic relations for testing scientific software: A machine learning approach using

- graph kernels, *Software Testing, Verification and Reliability* 26 (3) (2016) 245–269. doi:10.1002/stvr.1594.
- [19] F. Wang, X. Yang, M. Li, J. Liu, S. Yan, S. Fu, A dynamic identification method of metamorphic relation based on separation of input pattern and output pattern, in: *2022 9th International Conference on Wireless Communication and Sensor Networks (ICWCSN)*, ACM, 2022, pp. 16–24. doi:10.1145/3514105.3514109.
 - [20] S. Fu, X. Yang, M. Li, F. Wang, An identification method of image-based for output pattern of metamorphic relation in burnup calculation program, in: *2021 3rd International Academic Exchange Conference on Science and Technology Innovation (IAECST)*, IEEE, 2021, pp. 811–818. doi:10.1109/IAECST54258.2021.9695750.
 - [21] C. Tsigkanos, P. Rani, S. Müller, T. Kehrer, Variable discovery with large language models for metamorphic testing of scientific software, in: *Computational Science – ICCS 2023, Lecture Notes in Computer Science*, 2023, pp. 321–335. doi:10.1007/978-3-031-35995-8_23.
 - [22] H. Spieker, N. Belmecheri, A. Gotlieb, N. Lazaar, Evaluating human trajectory prediction with metamorphic testing, in: *Proceedings of the 9th ACM International Workshop on Metamorphic Testing (MET)*, 2024, pp. 34–40. doi:10.1145/3679006.3685071.
 - [23] H. Spieker, N. Lazaar, A. Gotlieb, N. Belmecheri, Metamorphic testing of multimodal human trajectory prediction, *Information and Software Technology* 188 (2025) 107890. doi:10.1016/j.infsof.2025.107890.
 - [24] D. J. Hiremath, M. Claus, W. Hasselbring, W. Rath, Towards automated metamorphic test identification for ocean system models, in: *2021 IEEE/ACM 6th International Workshop on Metamorphic Testing*, 2021, pp. 42–46. arXiv:2103.09782, doi:10.1109/MET52542.2021.00014.
 - [25] C. Mandrioli, S. Y. Shin, D. Bianculli, L. Briand, Testing cps with design assumptions-based metamorphic relations and genetic programming, *IEEE Transactions on Software Engineering* 51 (6) (2025) 1666–1684. doi:10.1109/TSE.2025.3563121.
 - [26] A. Duque-Torres, Selecting and constraining metamorphic relations, in: *Proceedings of the 2024 IEEE/ACM 46th International Conference*

- on Software Engineering: Companion Proceedings, 2024, pp. 212–216. doi:10.1145/3639478.3639781.
- [27] C.-a. Sun, S. Wu, X. Zhang, P.-L. Poon, Boosting metamorphic testing: A general metamorphic specification language and a supporting system, *ACM Transactions on Software Engineering and Methodology* (2026). doi:10.1145/3796225.
 - [28] M. Li, X. Yang, S. Yan, J. Liu, Y. Liu, J. Sun, A lightweight verification method based on metamorphic relation for nuclear power software, *Frontiers in Energy Research* 10 (2022) 788753. doi:10.3389/fenrg.2022.788753.
 - [29] M. Srinivasan, U. Kanewala, Metamorphic relation prioritization for effective regression testing, *Software Testing, Verification and Reliability* 32 (3) (2022) e1807. doi:10.1002/stvr.1807.
 - [30] P. Saha, U. Kanewala, Fault detection effectiveness of metamorphic relations developed for testing supervised classifiers, in: *2019 IEEE International Conference on Artificial Intelligence Testing (AITest)*, 2019, pp. 157–164. doi:10.1109/aitest.2019.00019.
 - [31] Y. Yang, T. F. M. Chui, Reliability assessment of machine learning models in hydrological predictions through metamorphic testing, *Water Resources Research* 57 (2021). doi:10.1029/2020WR029471.
 - [32] P. Reichert, K. Ma, M. Höge, F. Fenicia, M. Baity-Jesi, D. Feng, C. Shen, Metamorphic testing of machine learning and conceptual hydrologic models, *Hydrology and Earth System Sciences* 28 (11) (2024) 2505–2529. doi:10.5194/hess-28-2505-2024.
 - [33] H. F. Eniser, T. P. Gros, V. Wüstholtz, J. Hoffmann, M. Christakis, Metamorphic relations via relaxations: An approach to obtain oracles for action-policy testing, in: *Proceedings of the 31st ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*, 2022, pp. 52–63. doi:10.1145/3533767.3534392.
 - [34] M. Raissi, P. Perdikaris, G. E. Karniadakis, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, *Journal of Computational Physics* 378 (2019) 686–707. doi:10.1016/j.jcp.2018.10.045.

- [35] G. E. Karniadakis, I. G. Kevrekidis, L. Lu, P. Perdikaris, S. Wang, L. Yang, Physics-informed machine learning, *Nature Reviews Physics* 3 (2021) 422–440. doi:10.1038/s42254-021-00314-5.
- [36] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, A. Anandkumar, Fourier neural operator for parametric partial differential equations, in: *International Conference on Learning Representations*, 2021. arXiv:2010.08895, doi:10.48550/arXiv.2010.08895.
- [37] A. S. Krishnapriyan, A. Gholami, S. Zhe, R. M. Kirby, M. W. Mahoney, Characterizing possible failure modes in physics-informed neural networks, in: *Advances in Neural Information Processing Systems*, Vol. 34, 2021, pp. 26548–26560.
- [38] V. Gopakumar, A. Gray, L. Zanisi, T. Nunn, D. Giles, M. J. Kusner, S. Pamela, M. P. Deisenroth, Calibrated physics-informed uncertainty quantification, in: *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, Vol. 267 of *Proceedings of Machine Learning Research*, 2025. arXiv:2502.04406, doi:10.48550/arXiv.2502.04406.
- [39] S. Wang, Y. Teng, P. Perdikaris, Understanding and mitigating gradient flow pathologies in physics-informed neural networks, *SIAM Journal on Scientific Computing* 43 (5) (2021) A3055–A3081. doi:10.1137/20M1318043.
- [40] S. Baral, Y. Lee, S. Khanal, J. Jeon, Xrepit: A deep learning-computational fluid dynamics hybrid framework implemented in open-foam for fast, robust, and scalable unsteady simulations, *Computers & Fluids* 314 (2026) 107075. doi:10.1016/j.compfluid.2026.107075.
- [41] W. Wang, M. Hakimzadeh, H. Ruan, S. Goswami, Time-marching neural operator-finite element coupling: Ai-accelerated physics modeling, *Computer Methods in Applied Mechanics and Engineering* 446 (2025) 118319. doi:10.1016/j.cma.2025.118319.
- [42] H. Najafi, G. L. Lagerwall, J. Obeysekera, J. Liu, Machine learning in climate downscaling: A critical review of methodologies, physical consistency, and operational applications, *Water* 18 (2) (2026) 271. doi:10.3390/w18020271.

- [43] S. Yan, H. Zhu, Metamorphic testing on scientific programs for solving second-order elliptic differential equations, *Software Testing, Verification and Reliability* 35 (1) (2024). doi:10.1002/stvr.1912.
- [44] C. Stevens, K. Kjeer, R. M. Richard, E. F. Valeev, M. B. Cohen, Model assisted refinement of metamorphic relations for scientific software, in: 2025 IEEE/ACM 47th International Conference on Software Engineering: New Ideas and Emerging Results (ICSE-NIER), IEEE, 2025, pp. 31–35. doi:10.1109/ICSE-NIER66352.2025.00012.
- [45] X.-h. Yang, S.-y. Yan, J. Liu, M. Li, Hierarchical classification model for metamorphic relations of scientific computing programs, *Computer Science* 47 (11A) (2020) 557–561. doi:10.11896/jsjcx.200200015.
- [46] R. Verdecchia, E. Engström, P. Lago, P. Runeson, Q. Song, Threats to validity in software engineering research: A critical reflection, *Information and Software Technology* 164 (2023) 107329. doi:10.1016/j.infsof.2023.107329.
- [47] P. Ralph, et al., Empirical standards for software engineering research, *ACM SIGSOFT*, <https://github.com/acmsigsoft/EmpiricalStandards> (2021).